

Nearly Macro-free microKanren

Jason Hemann
Daniel P. Friedman
jason.hemann@shu.edu
dfried@cs.indiana.edu

Abstract

This paper describes changes to the microKanren implementation that make it more practical to use in a host language without macros. With the help of some modest runtime features common to most languages, we show how an implementer lacking macros can come closer to the expressive power that macros usually provide—with varying degrees of success. The result is a still functional microKanren that invites slightly shorter programs, and is relevant even to implementers that enjoy macro support. For those without it, we address some practical concerns that necessarily occur without macros so they can better weigh their options.

CCS Concepts: • Software and its engineering → Constraint and logic languages.

Keywords: logic programming, miniKanren, DSLs, embedding, macros

ACM Reference Format:

Jason Hemann and Daniel P. Friedman. 2023. Nearly Macro-free microKanren. In *Proceedings of Symposium on Trends in Functional Programming (TFP '23)*. ACM, New York, NY, USA, 7 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

The authors designed microKanren [5] as a compact relational programming language kernel to undergird a miniKanren implementation. Macros are used to implement the surrounding higher-level miniKanren operators and surface syntax. microKanren is often used as a tool for understanding the guts of a relational programming language through studying its implementation. By re-implementing miniKanren as separate surface syntax macros over a purely-function microKanren kernel, the authors hoped this separation would simultaneously aid implementers when studying the source

code, and also that the functional core would make the language easier to port to other functional hosts. To support both those efforts, they also chose to program in a deliberately small and workaday set of Scheme primitives.

The sum of those implementation restrictions, however, necessitates some awkward compromises in places including binary logical operators, one at a time local variable introduction, and leaks in the streams abstractions. These made the surface syntax macros seem practically mandatory, and fell short enough that we compromised on a purely functional kernel in a pedagogical exposition [4]. It also divided host languages into the macro language “haves” and macro-less “have nots”. Here, we bridge some of that divide by re-implementing parts of the kernel with some modest runtime features common to most languages.

In this paper we:

- show how to functionally implement more general logical operators, cleanly obviating some of the surface macros
- survey, the design space of purely functional implementation alternatives for the remaining macros in *The Reasoned Schemer, 2nd Ed's* [4] (*TRS2e*) core language implementation, and weigh the trade offs and real-world consequences
- suggest practical solutions for completely eliminating the macros in those places where the pure microKanren functional implementations had seemed impractical

This exercise resulted in some higher-level (variadic rather than just binary) operators, a more succinct kernel language, and enabled some performance improvement. Around half of the changes are applicable to any microKanren implementation, and the more concise goal combinators of Section 3 may also be of interest to implementers embedding goal-oriented languages like Icon. The other half are necessarily awkward yet practical strategies for those platforms lacking macro support. Our re-implementation’s source code for and the source for our experimental results is available at <https://github.com/jasonhemann/tfp-2023/>.

In Section 2, we illustrate by example what made surface syntax macros feel practically mandatory. In Section 3, we implement conjunction and disjunction, and in Section 4 we discuss the re-implementation of the impure operators. We discuss the remaining macros in Section 5. We close some

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

TFP '23, January 13-15, 2023, Boston, Massachusetts

© 2023 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/XXXXXXX.XXXXXXX>

outstanding questions on performance impacts of these implementation choices, and consider how Kanren language implementers outside of the Scheme family might benefit from these alternatives.

2 All Aboard!

We assume the reader is familiar with the miniKanren implementation of *The Reasoned Schemer, 2nd Ed (TRS2e)*. Although based on microKanren, this implementation makes some concessions to efficiency and safety and uses a few macros in the language kernel itself. In addition to that implementation, in this paper we make occasional references to earlier iterations such as Hemann et al. [6], an expanded archival version of the 2013 paper [5].

Haifa's Carmelit is the world's shortest subway system with a line of just six stations; an example sufficiently small that modeling it should be painless. But in microKanren, to model the order a passenger visit the stops riding that subway end to end requires 11 logical operator nodes, because microKanren only provides *binary* conjunctions and disjunctions. (Listing 11 contains this paper's alternative solution, requiring just three.) For a logic programming language, solely binary logical operators is too low level. To our eyes, this makes the superficial syntax macros practically mandatory, and host languages without a macro system are out of luck.

Moreover, the microKanren language doesn't offer the programmer much guidance in using that fine-grained control. For a series of n goals, the programmer can associate them to the left, to the right, or some mixtures of the two. The syntax does not obviously encourage any one choice. Subtle changes in program structure can have profound effects on performance, and mistakes are easy to make.

The TRS2e microKanren's soft-cut operator, `ifte`, is similarly low level. It permits a single test, a single consequent, and a single alternative. To build an if-then-else cascade, a microKanren programmer without the `conda` surface macro would need to code that unrolled cascade by hand.

The core TRS2e language implementation relies on macros `fresh`, `defrel`, and `run` to introduce new logic variables, globally define relations, and execute queries. Earlier implementations of those same behaviors via pure functional shallow embeddings, without macros, had some harsh consequences. We will revisit those earlier implementations and their trade-offs, survey the landscape of available choices, and suggest performant compromises for those truly without macros, thus increasing microKanren's *practical* portability.

3 `disj` and `conj` logical goal constructors

microKanren's binary `disj2` and `conj2` operators are goal combinators: they each take two goals, and produce a new goal. Disjunction and conjunction work slightly differently.

```
(define ((disj2 g1 g2) s)
  ($append (g1 s) (g2 s)))

(define ((conj2 g1 g2) s)
  ($append-map g2 (g1 s)))
```

Listing 1. microKanren `disj2` and `conj2`

```
(define ((disj . gs) s)
  (cond
    ((null? gs) (list))
    (else (D ((car gs) s) (cdr gs) s))))

(define (D ∞ gs s)
  (cond
    ((null? gs) ∞)
    (else
     (append ∞
              (D ((car gs) s) (cdr gs) s)))))
```

Listing 2. Eventual redefinition of `disj`

A *goal* is an outcome the program attempts to achieve, a goal can fail or succeed (and it can succeed many times). A goal executes with respect to a *state*, here the curried parameter `s`, and the result is a *stream* of states, and each entry is one state that results from achieving that goal in the given state. The `$append` function used in `disj2` is an internal kernel primitive that combines two streams into one, with an interleave mechanism to prevent starvation; the result is a stream of the ways to achieve the two goals' disjunction. The `$append-map` function used in `conj2` is to `$append` what the standard `append-map` is to `append`. The ways to achieve the conjunction of two goals are all the ways to achieve the second goal from the state after achieving the first goal. `$append-map` runs the second goal over the stream of results from the first goal, and combines the results of mapping together into a single stream representing the conjunction of the two, again with special attention to interleaving and starvation.

We want to implement disjunction and conjunction over arbitrary quantities of goals, as functions. These implementations should subsume the binary `disj2` and `conj2` and they also should not use `apply`. Further, they should not build any extraneous closures: unnecessarily building closures at runtime is always a bad idea. This re-implementation requires a host that supports variable arity functions, a widely available feature included in such languages as JavaScript, Ruby, Java, and Python. These languages do not generally support macros and hence are beneficiaries of this paper's contributions.

```

(define ((conj . gs) s)
  (cond
    ((null? gs) (list s))
    (else (C (cdr gs) ((car gs) s)))))

(define (C gs s∞)
  (cond
    ((null? gs) s∞)
    (else
     (C (cdr gs)
        (append-map∞ (car gs) s∞)))))

```

Listing 3. Eventual redefinition of conj

```

(define-syntax disj
  (syntax-rules ()
    ((disj g) g)
    ((disj g0 g1 g ...)
     (disj2 g0 (disj g1 g ...))))

(define (disj g . gs)
  (cond
    ((null? gs) g)
    (else (disj2 g (apply disj gs)))))

```

Listing 4. Deriving disj function from macro

Listings 2 and 3 show our new implementations. We re-implement these operators as shallow wrappers over simple folds. In each, the first steps are to dispense with the trivial case, and then to call a recursive help function that makes no use of variadic functions. That is, all of our focus will be on the recurring over the list *gs*. Unlike *D*, the function *C* does not take in the state *s*; the help procedure does not need the state for conjunction. In each recursive call, we accumulate by mapping (using that special delaying implementation of *append-map∞* for Kanren-language streams) the next goal in the list. This left-fold implementation of conjunction therefore left-associates the conjuncts.

3.1 Deriving semantic equivalents

A developer might derive these definitions as follows. We start with the definition of a recursive *disj* macro like one might define as surface syntax over the microKanren *disj₂*. As this is not part of the microKanren language itself, we would like to dispense with the macro and implement this behavior functionally. At the cost of an *apply*, we can build the corresponding explicitly recursive *disj* function. Since *disj* produces and consumes goals, we can η -expand the definition in Listing 4 by a curried parameter *s*. We then split *disj* into two mutually-recursive procedures, to build the variant in Listing 5.

```

(define ((disj g . gs) s)
  (D g gs s))

(define (D g gs s)
  (cond
    ((null? gs) (g s))
    (else ((disj2 g (apply disj gs)) s))))

```

Listing 5. An η -expanded and split definition of disj

```

(define (D g gs s)
  (cond
    ((null? gs) (g s))
    (else
     (append∞ (g s)
              (D (car gs) (cdr gs) s)))))

```

Listing 6. Derivation of disj function definition

```

(define-syntax conj
  (syntax-rules ()
    ((conj g) g)
    ((conj g g1 gs ...)
     (conj (conj2 g g1) gs ...)))

```

```

(define (conj g . gs)
  (cond
    ((null? gs) g)
    (else
     (apply conj
              (cons (conj2 g (car gs)) (cdr gs)))))

```

Listing 7. conj₂-based conj macro and function

We can replace the call to *disj₂* in Listing 5 by its definition in terms of *append∞* and perform a trivial β -reduction. The explicit *s* argument suggests removing the call to *apply* and making *D* recursive. The result is the version of *D* in Listing 6. The definition of *disj* remains unchanged from Listing 5. In both clauses of *D* we combine *g* and *s*, this suggests constructing that stream in *disj* and passing it along. Adding the trivial base case to that *disj* yields the definition in Listing 2.

We can derive the definition of *conj* from Listing 3 via a similar process. Starting with the variadic function based on the macro in Listing 7, we first η -expand and split the definition.

We next substitute for the definitions of *conj* and *conj₂*.

Finally, since *C* only needs *s* to *build* the stream, we can assemble the stream on the way in—instead of passing in *g* and *s* separately, we pass in their combination as a stream. The function is tail recursive, we can change the signature

```
(define ((conj g . gs) s)
  (C g gs s))

(define (C g gs s)
  (cond
    ((null? gs) (g s))
    (else
     ((apply conj
              (cons (conj2 g (car gs)) (cdr gs)))
      s))))
```

Listing 8. Derivation of split conj function definition

```
(define (C g gs s)
  (cond
    ((null? gs) (g s))
    (else
     (C (λ (s) (append-map∞ (car gs) (g s)))
        (cdr gs)
        s))))
```

Listing 9. Replacing apply in C function definition

in the one and only external call and the recursive call. The result, after adding the trivial base case to conj, is shown in Listing 3.

Both the functional and the macro based versions of Listing 7 use a left fold over the goals, whereas the versions of disj use a right fold. This is not an accident. Folklore suggests that left associating conjunctions tends to improve the performance of miniKanren’s interleaving search. The authors know of no thorough algorithmic proof of such claims, but see for instance discussions and implementation in [9] for some of the related work so far. In ??, we display the results of some micro benchmarks that suggest the same. We have generally, however, resorted to small step visualizations of the search tree to explain the performance impact. The authors believe it is worth considering if we can make an equally compelling argument for this preference through equational reasoning and comparing the implementations of functions.

The benefits of a left-fold over conjunctions becomes a little more obvious by comparison to a right-fold implementation after we η -expand, unfold to a recursive help function, substitute in the definition of conj₂, and β -reduce. From there, we cannot (easily) replace the apply call by a recursive call to C, because we are still waiting for an s. We can only abstract over s and wait; we show the upshot of this sequence in Listing 10. Since we know that any call to append-map∞ we construct will always yield a result, the version in Listing 3 is tail recursive. The equivalent right-fold

```
(define ((conj g . gs) s)
  (C gs (g s)))

(define (C gs s∞)
  (cond
    ((null? gs) s∞)
    (else
     (append-map∞
      (λ (s) (C (cdr gs) ((car gs) s)))
      s∞))))
```

Listing 10. A right-fold variant of conj after some derivations

implementation needs to somehow construct a closure for every recursive call. Basic programming horse sense suggests the more elegant variant from Listing 3.

The new disj and conj functions are, we believe, sufficiently high-level for programmers in implementations without macros. Though this note mainly concerns working towards an internal macro-less kernel language, it may also have something to say about the miniKanren-level surface syntax, namely that even the miniKanren language could do without its conde syntax (a disjunction of conjunctions that looks superficially like Scheme’s cond) and have the programmer use these new underlying logical primitives. We implement carmelit-subway as an example in Listing 11, and it looks far cry better than the 11 binary logical operator nodes the programmer would have needed to write in a microKanren language without macros.

```
(defrel (carmelit-subway a b c d e f)
  (disj
    (conj (== a 'carmel-center)
          (== b 'golomb)
          (== c 'masada)
          (== d 'haneviim)
          (== e 'hadar-city-hall)
          (== f 'downtown))
    (conj (== a 'downtown)
          (== b 'hadar-city-hall)
          (== c 'haneviim)
          (== d 'masada)
          (== e 'golomb)
          (== f 'carmel-center))))
```

Listing 11. A new Carmelit subway without conde

4 Tidying up the Impure Operators

The miniKanren conda operator that provides nested “if-then-else” behavior relies on the microKanren ifte underlying it. The definition of conda (see Listing 12) requires one or more conjuncts per clause and one or more clauses. The

```
(define-syntax conda
  (syntax-rules ()
    ((conda (g0 g ...) (conj g0 g ...))
     ((conda (g0 g ...) ln ...)
      (ifte g0 (conj g ...) (conda ln ...)))))

(define ((ifte g1 g2 g3) s)
  (let loop ((s∞ (g1 s)))
    (cond
      ((null? s∞) (g3 s))
      ((pair? s∞)
       (append-map∞ g2 s∞))
      (else (lambda ()
               (loop (s∞)))))))
```

Listing 12. A typical implementation of conda

```
(define ((conda q a . q-and-a*) s)
  (A (q s) a q-and-a* s))

(define (A s∞ a q-and-a* s)
  (cond
    ((null? s∞)
     (cond
       ((null? (cdr q-and-a*)) ((car q-and-a*) s))
       (else (A ((car q-and-a*) s)
                 (cadr q-and-a*)
                 (cddr q-and-a*)
                 s))))
    ((pair? s∞) (append-map∞ a s∞))
    (else (lambda () (A (s∞) a q-and-a* s)))))
```

Listing 13. A functional conda implementation

last line of conda contains the only place in the implementation that relies structurally on permitting nullary conjunctions, or disjunctions, of goals. Everywhere else conjunctions are one-or-more, and this one structural dependency is off-putting. Having already broached the topic of changing surface syntax, we mention a temptation to rewrite the second pattern in miniKanren’s conda to demand *two* or more goals in each if-then clause and removing the dependency.

Some microKanren programmers without macros would be perfectly satisfied just using ifte directly, especially so given the research community’s focus on purely relational programming. But just as the standard forked if begat McCarthy’s if notation and cond, a programmer may eventually feel the need for a nested implementation. Here is a functional implementation of that cascade behavior.

The implementation in Listing 13 includes the delay-and-restart behavior of ifte together with conda’s logical cascade. The s_∞ can be either empty, non-empty, or a function of no arguments. In the last case, we invoke s_∞. Rather than

```
(define (once g)
  (lambda (s)
    (let loop ((s-inf (g s)))
      (cond
        ((null? s-inf) '())
        ((pair? s-inf)
         (cons (car s-inf) '()))
        (else (lambda ()
                 (loop (s-inf))))))))
```

Listing 14. The once function

```
(define (subtleo x)
  (Zzz
   (disj
    (subtleo x)
    (= x 'cat))))
```

Listing 15. Omitting the delay is a subtle bug

building a largely redundant implementation of conda, we expose the higher-order goal once to the user. The definition of once in Listing 14 is taken directly from [4]. The programmer can simulate conda by wrapping once around every test goal.

5 Remainders and Practicalities

We have not fully obviated the use of macros. In this section we collect together some workarounds to obviate macros in the rest of the implementation. Some of these come with significant drawbacks. With these, however, a programmer in even a pedestrian functional language should be able to directly translate the implementation and our test programs.

define. The microKanren programmer can just use their host language’s define feature to construct relations as host-language functions, and manually introduce the delays in relations. This may be a larger concession than it looks, since it exposes the delay and interleave mechanism to the user, and both correct interleaving and even the termination of relation *definitions* rely on a whole-program correctness property of relation definitions having a delay. Listing 15 relies on a help function Zzz to introduce delays, akin to some earlier implementations [5]. Another downside worth mentioning is the programmer must now take extra care not to provide multiple goals to define. The define form will treat all but the last expression as statements and silently drop them, rather than conjoin them as in defrel. That small *gotcha* can be subtle but significant drawback.

fresh. In any implementation there must be some mechanism to produce the next fresh variable. For example, we treat the natural numbers as an indexed set of variables, and we thread the current index through the computation. We

```
(call/initial-state 1
  (let ((q (var 'q)))
    (conj
      (let ((x (var 'x)))
        (== q x))
      (reify q))))
```

Listing 16. Queries as expressed with global-state variables

use `add1` to get the next index; to go from index to variable is the identity function. For another example, we could represent each variable using a unique memory location, *sokuza-kanren* [8] style, and the operation to produce a new variable requires introducing an unused memory location. Depending on the implementation of variables, you may also need additional functions to support your implementation of variables. If variables are not from an indexed set, you may also need an operation to (re) construct specifically the first element of the set, or otherwise store that value for later re-use.

Of course, one of these approaches requires memory allocation and external global state, while the other does not. Furthermore, the latter approach models logic variables as coming from a single global pool rather than reusing them separately across each disjunct, and so requires some global store and strictly more logic variables overall.

With this latter approach, however, we can expose `var` directly to the programmer and the programmer can use `let` bindings to introduce several logic variables simultaneously.

run. We note that we can also implement `run` without using macros. Using the purely-functional implementation of logic variables, the definitions of `run` and `run*` easily translate to functions like `call/initial-state` [5]. The query is itself expressed as a goal that introduces the first logic variable `q`. The pointer-based logic variable approach forces the programmer to explicitly invoke `reify` as though it were a goal as the last step of executing the query, as in Listing 16.

6 Future Work

This note shows how to provide a somewhat more concise core language that significantly reduces the need for macros. The result almost rivals the expressivity of the full “microKanren + macros” approach. Variadic functions make this implementation much more convenient for the end programmer, and Scheme’s polyvariadic function syntax ensures at a host-language level that the microKanren programmer provides at least one parameter to `conj` and `disj`.

The old desugaring macros do not seem to suggest how to associate the calls to the binary primitives—both left and right look equally nice. Forcing ourselves to program the

solution functionally, and the restrictions we placed on ourselves in this reimplementing, removed a degree of implementation freedom and led us to what seems like the right solution.

The result is closer to the design of Prolog, where the user represents conjunction of goals in the body of a clause with a comma and disjunction, either implicitly in listing various clauses or explicitly with a semicolon. We assume it is agreed that our definitions of `disj` and `conj` themselves are sufficiently high-level operators for a surface language and that the zero-element base cases are at best unnecessary and likely undesirable; given the opportunity to define a surface language and its desugaring, we really shouldn’t tempt the programmer by making undesirable programs representable when we can avoid it.

Techniques for implementing `defrel`, `fresh` and `run` (and `run*`) without macros come with serious drawbacks. These include exposing the implementation of streams and delays, and the inefficiency and clumsiness of introducing variables one at a time, or the need to reason with global state.

From time to time we find that the usual miniKanren implementation is *itself* lower-level than we would like to program with relations. Early microKanren implementations restrict themselves to syntax-rules macros. Some programmers use macros to extend the language further as with *matche* [7]. Some constructions over miniKanren, such as *minikanren-ee* [1], may rely on more expressive macro systems like *syntax-parse* [2].

We would still like to know if our desiderata here are *causally* related to good miniKanren performance. Can we reason at the implementation level and peer through to the implications for performance? If left associating `conj` is indeed uniformly a dramatic improvement, the community might consider reclassifying left-associative conjunction as a matter of correctness rather than an optimization, as in “tail call optimization” vs. “Properly Implemented Tail Call Handling” [3]. Regardless, we hope this document helps narrow the gap between implementations in functional host languages with and without macro systems and helps implementers build more elegant, expressive and efficient Kanrens in their chosen host languages.

Acknowledgments

Thanks to Ken Shan and Jeremy Siek, for helpful discussions and debates during design decision deliberations. Thanks also to Greg Rosenblatt and Michael Ballantyne for their insights and suggestions. We would also like to thank our anonymous reviewers for their insightful contributions.

References

- [1] Michael Ballantyne, Alexis King, and Matthias Felleisen. “Macros for domain-specific languages.” In: *Proceedings of the ACM on Programming Languages* 4.OOPSLA (2020), pp. 1–29.
- [2] Ryan Culpepper. “Fortifying macros.” In: *Journal of functional programming* 22.4-5 (2012), pp. 439–476.
- [3] Matthias Felleisen. *Re: Question about tail recursion*. 2014. URL: <https://lists.racket-lang.org/users/archive/2014-August/063844.html>.
- [4] Daniel P. Friedman et al. *The Reasoned Schemer, Second Edition*. The MIT Press, Mar. 2018. ISBN: 0-262-53551-3. URL: mitpress.mit.edu/books/reasoned-schemer-0.
- [5] Jason Hemann and Daniel P. Friedman. “μkanren: A Minimal Functional Core for Relational Programming.” In: *Scheme 13*. 2013. URL: <http://schemeworkshop.org/2013/papers/HemannMuKanren2013.pdf>.
- [6] Jason Hemann et al. “A Small Embedding of Logic Programming with a Simple Complete Search.” In: *Proceedings of DLS '16*. ACM, 2016. URL: <http://dx.doi.org/10.1145/2989225.2989230>.
- [7] Andrew W Keep et al. “A pattern matcher for miniKanren or How to get into trouble with CPS macros.” In: *Technical Report CPSLO-CSC-09-03* (2009), p. 37.
- [8] Oleg Kiselyov. *The taste of logic programming*. 2006. URL: <http://okmij.org/ftp/Scheme/misc.html#sokuza-kanren>.
- [9] Gregory Rosenblatt et al. “First-order miniKanren representation: Great for tooling and search.” In: *Proceedings of the miniKanren Workshop* (2019), p. 16.