

# Some criteria for implementations of conjunction and disjunction in microKanren

Anonymous Author(s)

## Abstract

We describe here how, at the cost of an additional language feature—namely variadic and “rest argument” functions—we can implement a slightly higher-level kernel language. We implement conjunction and disjunction functions that preserve search order and that do not create superfluous closures. The result is a shorter and simpler overall implementation that is more elegant.

## ACM Reference Format:

Anonymous Author(s). 2018. Some criteria for implementations of conjunction and disjunction in microKanren. In *Proceedings of miniKanren Workshop (mKW '22)*. ACM, New York, NY, USA, 7 pages. <https://doi.org/XXXXXXX.XXXXXXX>

## 1 Introduction

microKanren [3] is a compact approach to implementing a relational programming language. The microKanren approach has worked out well as a tool for understanding the guts of a relational programming language through studying its implementation. The microKanren reimplemention separates surface syntax macros from function definitions. In doing so, the authors hoped this separation will simultaneously aid future would-be implementers when studying the source code, and also that the functional core would make the language easier to port to other functional host languages. To support both those efforts, they also chose to program in a deliberately small and workaday set of Scheme primitives.

This note shows how an implementer using a host language with variadic (any number of arguments) functions can build a somewhat more powerful kernel language and obviate macros for conjunction and disjunction. Until now microKanren implementations have divided the programming languages of the world more or less into the expressive macro system “haves” and the macro system-less “have nots”; variadic functions permit an intermediate point in

the language design space. Even implementers in languages *with* sufficiently macro systems might choose to replace existing less powerful language primitives by our comparatively short but more powerful variants. Beyond that, this more powerful kernel language obviates several macros in the desugaring layer, and eventually leads us to see merit in modifying the core language’s syntax. This is, of course, subjective.

Beyond being a conceptually simpler foundation for a full Kanren language at the level of *The Reasoned Schemer, Second Edition*, this implementation appears to be more *efficient*. Right-associative conjunction is already conjectured in miniKanren implementation folklore to be an inefficient design choice. We were pleasantly surprised that thinking about the expressivity of primitives and code improvement “horse sense” led us to this same idea. That very horse sense might lead one to suggest these improved variants should actually be the external programming constructs. This, too, is very speculative.

In Section 2, we briefly revisit microKanren implementations and illustrate why surface syntax macros had seemed practically mandatory. In Section 3, we implement conjunction and disjunction, and in Section 4 we discuss the reimplementation of the impure operators. We close with a discussion of some performance impacts of these implementation choices, and consider how Kanren language implementers outside of the Scheme family might benefit from these alternative implementations.

## 2 All Aboard!

We assume the reader is familiar with the miniKanren language as described in *The Reasoned Schemer, Second Edition* and in particular with microKanren implementations. See [minikanren.org](http://minikanren.org) for many implementations across multiple host languages.

The world’s shortest subway system is Haifa’s Carmelit, with only six stations. The system is a line, so its trains travel back and forth. If we wanted to describe the order in which we rode the train from the beginning to the end, we could use the `carmelit-subway` relation in Listing 1. We express this relation using the more compact miniKanren syntax because unfolding this relation into binary conjunctions and disjunctions would be painful. We certainly *could* write it out by hand—in fact, we could write it out many ways. We could nest those conjunctions to the left, or to the

---

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

mKW '22, September 15, 2022, Ljubljana, Slovenia

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/XXXXXXX.XXXXXXX>

```

111 (defrel (carmelit-subway a b c d e f)
112   (conde
113     ((= a 'carmel-center)
114      (= b 'golomb)
115      (= c 'masada)
116      (= d 'haneviim)
117      (= e 'hadar-city-hall)
118      (= f 'downtown))
119     ((= a 'downtown)
120      (= b 'hadar-city-hall)
121      (= c 'haneviim)
122      (= d 'masada)
123      (= e 'golomb)
124      (= f 'carmel-center))))

```

**Listing 1.** A miniKanren implementation of the Carmelit subway.

right, or try and reduce the indentation by trying to keep them somewhat balanced: the program itself does not seem to obviously encourage one particular choice.

In a microKanren in a macro-less host language, a programmer could get bothered by low-level details such as stacking binary conjunction and disjunction. A small programming mistake, such as incorrectly nesting a conjunct somewhere in a long conjunction chain, can cause relatively obscure errors and reading those error messages forces the programmer out of the microKanren-DSL thinking, and back into the host programming language’s debug model.

Listing 2 shows typical macro-based implementations of `conde` and the underlying conjunction and disjunction operations. Here, `#s` and `#u` represent primitive goals that unconditionally succeed and fail, respectively.

### 3 conj and disj logical goal constructors

The `conj` and `disj` macros of Listing 2 are not *quite* defined as simple recursions over their binary functional primitives. The definition of 0-way conjunction (disjunction) is independent of the unfolding of `conj2` (`disj2`). In a sense the `conj` and `disj` macros confuse and entangle primitive success and failure goals with those recursive unfoldings in terms of the binary operators.

Those zero-way logical operation base cases don’t add much. The programmer who tries to write an elegant, efficient solution to a pure relational programming task would not encounter these additional base cases in his `conde` expressions. The programmer knows statically how these goals should behave, so there is no benefit to executing them. A conjunction of no goals would simply succeed, and a disjunction of no goals would simply fail. Truthfully, only the impure `conda` operator seems to require these additional base cases (discussed further in Section 4) they are for all intents and purposes superfluous.

```

166 (define-syntax conde
167   (syntax-rules ()
168     ((conde (g ...) ...)
169      (disj (conj g ...) ...))))
170
171 (define-syntax conj
172   (syntax-rules ()
173     ((conj) #s)
174     ((conj g) g)
175     ((conj g0 g ...) (conj2 g0 (conj g ...))))
176
177 (define ((conj2 g1 g2) s)
178   (append-map g2 (g1 s)))
179
180 (define-syntax disj
181   (syntax-rules ()
182     ((disj) #u)
183     ((disj g) g)
184     ((disj g0 g ...) (disj2 g0 (disj g ...))))
185
186 (define ((disj2 g1 g2) s)
187   (append (g1 s) (g2 s)))
188

```

**Listing 2.** Macro based implementations of `conj` and `disj`.

Here, we want to implement conjunction and disjunction over one-or-more goals as functions. These implementations will rely on variadic function syntax, but should not require `apply` or rely on the binary `conj2` and `disj2`. Further, they should not build any extraneous closures: unnecessarily building closures at runtime is always a bad idea.

Listing 3 shows our new implementation. Mandating one-or-more arguments lets us re-implement these operators as shallow wrappers over simple folds. The first step in each is merely to remove the rest argument `gs` and act as if there were no need for a rest argument. That is, all of our focus will be on the *list*, `gs`. The function `C` does not take in the state `s`; the help procedure does not need the state for conjunction. In each recursive call, we accumulate by mapping (using the special delaying implementation of `append-map` for miniKanren streams) the next goal in the list. This left-fold implementation of conjunction therefore left-associates the conjuncts.

#### 3.1 Implementation correctness

Equipped with variadic functions, a developer might discover these definitions as follows. We start with the definition of `disj` from Listing 3. At the cost of an `apply`, we can define this as a single explicitly recursive function.

```

216 (define ((disj g . gs) s)
217   (cond
218     ((null? gs) (g s))
219     (else

```

```

221 (define ((conj g . gs) s)
222   (C gs (g s)))
223
224 (define (C gs s∞)
225   (cond
226     ((null? gs) s∞)
227     (else
228      (C (cdr gs)
229         (append-map∞ (car gs) s∞))))))
230
231 (define ((disj g . gs) s)
232   (D g gs s))
233
234 (define (D g gs s)
235   (cond
236     ((null? gs) (g s))
237     (else
238      (append∞ (g s)
239                 (D (car gs) (cdr gs) s))))))
240
241 (define-syntax conde
242   (syntax-rules ()
243     ((conde (g g1 ...) ...)
244      (disj (conj g g1 ...) ...))))

```

### Listing 3. Re-implementations of conj and disj.

```

247
248 (define ((disj g . gs) s)
249   (cond
250     ((null? gs) (g s))
251     (else ((λ (s)
252              (append∞ (g s)
253                        ((apply disj gs) s)))
254            s))))
255
256   (append∞ (g s)
257             ((apply disj gs) s))))

```

We inverse  $\beta$ -substitute in the recursive case, to draw out a similarity between the two clauses.

We can replace the call to `append∞` by the definition of `disjz`.

```

261 (define ((disj g . gs) s)
262   (cond
263     ((null? gs) (g s))
264     (else ((disjz g (apply disj gs)) s))))
265

```

Now `s` is a curried parameter we can  $\eta$  away. We provide the completed functional definition and a comparable to the macro-based implementation in Listing 4.

We can demonstrate the correctness of our `conj` from Listing 3 in the same manner. We first deconstruct the stream accumulator back into its two constituent pieces.

```

272 (define ((conj g . gs) s)
273   (C g gs s))
274
275

```

```

276 (define (disj g . gs)
277   (cond
278     ((null? gs) g)
279     (else (disjz g (apply disj gs)))))
280
281 (define-syntax disj
282   (syntax-rules ()
283     ((disj g) g)
284     ((disj g gs ...) (disjz g (disj gs ...)))))
285

```

### Listing 4. Derived disj<sub>z</sub>-based function and macro.

```

289 (define (C g gs s)
290   (cond
291     ((null? gs) (g s))
292     (else
293      (C (λ (s) (append-map∞ (car gs) (g s)))
294         (cdr gs)
295         s))))
296

```

This we can collapse into a single recursive function, bringing the goal to the front.

```

299 (define ((conj g . gs) s)
300   (cond
301     ((null? gs) (g s))
302     (else
303      ((apply conj
304              (cons
305                (λ (s) (append-map∞ (car gs) (g s)))
306                (cdr gs))
307              s)))))
308

```

We can replace the abstraction over `append-map∞`, which is to say the body of `conjz`, by its definition.

```

311 (define ((conj g . gs) s)
312   (cond
313     ((null? gs) (g s))
314     (else
315      ((apply conj
316              (cons (conjz g (car gs)) (cdr gs))
317              s)))))
318

```

By  $\eta$  reducing over the definition, we produce a simple variadic function. We show the result, and a comparable macro-based definition in Listing 5.

Both the variadic function based and the macro based versions of Listing 5 use a left fold over the goals, whereas the versions of `disj` use a right fold. This is not an accident, and we believe this is *worth further consideration*. It is miniKanren folklore that left associating conjunctions tends to improve miniKanren performance, but we have generally resorted to small step visualizations of the search tree to demonstrate why that might be. We can make an equally compelling argument for this preference through equational

```

331 (define (conj g . gs)
332   (cond
333     ((null? gs) g)
334     (else
335      (apply conj
336       (cons (conj2 g (car gs)) (cdr gs))))))

```

```

337
338 (define-syntax conj
339   (syntax-rules ()
340     ((conj g) g)
341     ((conj g g1 gs ...)
342      (conj (conj2 g g1) gs ...)))

```

**Listing 5.** Derived conj<sub>2</sub>-based function and macro.

```

345
346 (define (conj g . gs)
347   (cond
348     ((null? gs) g)
349     (else (conj2 g (apply conj gs)))))

```

**Listing 6.** A right-fold variant of conj.

reasoning. Listing 6 shows a right-fold variant of conj. Looking at either the functional or macro definitions built from conj<sub>2</sub>, it's difficult to assert any preference for one over the other. The apply primitive and the macros both get in the way.

The choice becomes a little more obvious after we η-expand unfold to a recursive help function, substitute in the definition of conj<sub>2</sub>, and β-reduce.

```

361 (define ((conj g . gs) s)
362   (C g gs s))
363
364 (define (C g gs s)
365   (cond
366     ((null? gs) (g s))
367     (else (append-map-∞ (apply C gs) (g s)))))
368
369 (define ((conj g . gs) s)
370   (C gs (g s)))

```

Here, we cannot (easily) replace the apply call by a recursive call to C, because we are still waiting for an s. We can only abstract over s and wait.

```

371
372 (define (C g gs s)
373   (cond
374     ((null? gs) (g s))
375     (else (append-map-∞ (apply C gs) (g s)))))
376
377
378
379 (define ((conj g . gs) s)
380   (C gs (g s)))
381
382 (define (C gs s∞)
383   (cond
384     ((null? gs) s∞)

```

```

385
386 (else
387   (append-map-∞
388     (λ (s) (let ((t∞ ((car gs) s)))
389              (C (cdr gs) t∞)))
390     s∞))))

```

Since we know that any call to append-map<sup>∞</sup> we construct will always yield a result, the other one is tail recursive. This version is not, and it builds closures. We can still reconstruct the stream argument, and come up with an approximately analogous right-fold version.

```

391
392
393
394
395
396
397 (define ((conj g . gs) s)
398   (C gs (g s)))
399
400
401 (define (C gs s∞)
402   (cond
403     ((null? gs) s∞)
404     (else
405      (append-map-∞
406        (λ (s) (C (cdr gs) ((car gs) s)))
407        s∞))))

```

So the equivalent right-fold implementation needs to construct a closure for every recursive call. If we want to implement a variadic version that does not rely on a primitive conj<sub>2</sub> and does not resort to apply, we have the two aforementioned choices. Basic programming horse sense suggests the variant from Listing 3.

Left associating conjuncts is so important that it dramatically improves miniKanren's search. A full, thorough, explanation of the behavior is outside the scope of our work here, but see discussions by Rosenblatt and Ballantyne [[here, nowhere](#)].

DONE UP UNTIL HERE

miniKanren's biased search heuristic causes some operational differences between this reimplement and the older one. With this implementation, instead of the *first* goal being the most heavily weighted in the search, here the *last* goal is most heavily weighted. Because *n*-ary disjunction is now built in to the functional definition, this behavior is now transparently visible in the implementation of D. Imagine for a moment that in the given state *s* each goal in (g<sub>1</sub> . gs) produces infinitely many answers. The append is doing a fair interleaving between its two stream arguments: giving one a little time to work and then handing off to the other. At each recursion, approximately half the answers in that stream come from the most recently evaluated goal, and so the resources dedicated to all the preceding goals is halved, in order to accommodate this new one.

Declarative programmers oughtn't need concern themselves with the operational search behavior, and since miniKanren provides a complete search, a query should still return the same bag of answers, so we do not see this as a problem; we just consider it an interesting distinction. If however operationally-aware programmers found the "reversed

```

441 (defrel (carmelit-subway a b c d e f)
442   (disj
443     (conj (== a 'carmel-center)
444           (== b 'golomb)
445           (== c 'masada)
446           (== d 'haneviim)
447           (== e 'hadar-city-hall)
448           (== f 'downtown))
449     (conj (== a 'downtown)
450           (== b 'hadar-city-hall)
451           (== c 'haneviim)
452           (== d 'masada)
453           (== e 'golomb)
454           (== f 'carmel-center))))

```

**Listing 7.** A reimplemented Carmelit subway without conde.

weighting” unintuitive, some language front end could rearrange the programmer’s arguments. For example, reversing a list of arguments is an exemplary use of an APS macro.

IS IT ACTUALLY SO? HOW WOULD I IMPLEMENT FOLD■  
THE OTHER WAY.

Though this note mainly concerns the choice to implement surface language functionality as macros, it may also point to these as more natural user-level primitives than conde. An implementation could choose to forego conde and provide just those underlying logical primitives conj and disj to the user.

## 4 Cleaning up the Impure Operators

The zero-ary conjunction and disjunction base cases both forced primitive goals #s and #u into the language and also hid the opportunity of functional left-fold implementations as helpers to variadic surface functions. The zero-ary conjunction case was itself implemented as an answer to an earlier problem in implementing an if-then-else soft-cut operator. The zero-ary disjunction base case was not *really* required, but given that we needed one the other came along for the ride, so to speak.

The operators conda and conde look superficially similar, syntactically. Semantically though, the disjunction and nested “if-then-else” behaviors are quite different, and implementing the desired behavior for this conda from existing combinators raises some oddities. The definition of conda already rules out zero-way conjunctions of goals in the clause body as well as zero-way disjunctions of such clauses. Another oddity is that, for a conda syntax that permits one-or-more-goals in each clause, the final clause of the conda definition is the one and only place in the implementation that may require a nullary conjunction of goals.

There are work-arounds. One could:

```

496 (define-syntax conda
497   (syntax-rules ()
498     ((conda (g0 g ...) (conj g0 g ...))
499       ((conda (g0 g ...) ln ...)
500         (ifte g0 (conj g ...) (conda ln ...))))))
501
502 (define ((ifte g1 g2 g3) s)
503   (let loop ((s∞ (g1 s)))
504     (cond
505       ((null? s∞) (g3 s))
506       ((pair? s∞)
507        (append-map g2 s∞))
508       (else (lambda ()
509                (loop (s∞)))))))
510

```

**Listing 8.** A typical implementation of conda.

1. Special case one-goal clauses in the last clause of the conda implementation with a special ifte variant.
2. Special case that last clause by introducing #s as an additional goal.
3. Mandate that all clauses in conda contain at least two goals.
4. Mandate that all clauses *except the final default clause* contain at least two goals.
5. Build a special-case conjunction solely for this purpose.

In all but the first of these choices, either the implementation implicitly or the syntax explicitly can force additional unneeded executions of unwanted goals. The first choices seems even worse for introducing another complex function definition that is nearly identical to one used in the below case.

Here too, removing macros and relying on variable arity function definition provides a more elegant solution. Some programmers would be perfectly satisfied to have lived forever with ifte itself. But just as the standard forked if be-gat McCarthy’s if notation and cond, we can suspect that a programmer would eventually feel the need for a nested implementation.

The implementation in Listing 9 includes the delay-and-restart behavior of ifte together with conda’s logical cascade.

### 4.1 condu

We have not yet addressed condu. We could implement this for functional programmers with a variant of Listing 9. Instead, we choose to make once a primitive higher-order goal. This relies on a local let-bound named recursive procedure.

## 5 Conclusion

narrows the gap between implementations in functional host■  
languages with and without macro systems.



```

551 (define-syntax conda
552   (syntax-rules ()
553     ((conda (g0 g ...) (conj g0 g ...))
554      ((conda (g0) ln ...)
555       (if g0 (conda ln ...)))
556      ((conda (g0 g ...) ln ...)
557       (ifte g0 (conj g ...) (conda ln ...)))))
558
559 (define ((if g1 g3) s)
560   (let loop ((s∞ (g1 s)))
561     (cond
562       ((null? s∞) (g3 s))
563       ((pair? s∞) s∞)
564       (else (lambda ()
565                (loop (s∞)))))))
566
567 (define ((ifte g1 g2 g3) s)
568   ...)

```

**Listing 9.** An inauspicious re-implementation of `conda`.

```

571
572 (define ((conda q a . q-and-a*) s)
573   (A (q s) a q-and-a* s))
574
575 (define (A s∞ a q-and-a* s)
576   (cond
577     ((null? s∞)
578      (cond
579        ((null? (cdr q-and-a*)) ((car q-and-a*) s))
580        (else (A ((car q-and-a*) s)
581                  (cadr q-and-a*)
582                  (cddr q-and-a*) s))))))
583   ((pair? s∞) (append-map∞ a s∞))
584   (else (lambda () (A (s∞) a q-and-a* s))))))
585

```

**Listing 10.** A functional `conda` implementation.

```

586
587
588 (define (once g)
589   (lambda (s)
590     (let loop ((s∞ (g s)))
591       (cond
592         ((null? s∞) '())
593         ((pair? s∞)
594          (cons (car s∞) '()))
595       (else (lambda ()
596                (loop (s∞)))))))
597

```

**Listing 11.** A functional `once` implementation.

More languages than ever before support variable arity functions/methods (aka `varargs` aka `slurpy methods`), including Raku, Java, and Ruby, to name just a few. Many of these languages also support polyvariadic functions. The

authors hope that this document helps implementers build more elegant, expressive and efficient Kanrens in their chosen host languages.

Languages with variadic but without polyvariadic (that is to say “at least  $k$  arguments”) functions—how do they fare here?

This is closer to the design of Prolog, where the user represents conjunction of goals in the body of a clause with a comma and disjunction, either implicitly in listing various clauses or explicitly with a semicolon.

Languages without expressive macro systems would still have to resort to exposing implementations of streams and *some* lower level operations like introducing variables one at a time. But!

We take it as granted that conjunction and disjunction themselves are sufficiently high-level operators for a surface language and that the 0-element base cases are at best unnecessary and likely undesirable.

Given the opportunity to define a surface language and its desugaring, we really shouldn’t tempt the programmer by making undesirable programs representable when we can avoid doing so.

Of course, no implementer *needs* a macro system to implement a shallow embedding of an LP language. In our implementations, we still use macros to implement `defrel` and `fresh`, the former to prevent exposing the implementation of streams and delays, and the latter both for the added efficiency and to avoid the awkwardness of introducing each variable one at a time.

From time to time we find that the usual miniKanren implementation is itself lower-level than we would like to program with relations. Expert miniKanren programmers use macros to extend the language yet again, as with `matche`, and also with wholly more expressive and highly optimized language forms, as in the original Kanren. It is nice to find we can spread some of that extra expressiveness to the wider Kanren language community.

Even implementers using and programmers studying implementations written in languages *with* sufficiently expressive macro systems can benefit from these improvements. There are benefits to using a limited set of features in a language. Based on an evaluation of `minikanren.org`, earlier microKanren implementers measured their results by the number of lines in the core implementation, the number of additional lines for the surface syntax, and the limited features they required from the host. So replacing

Early microKanren implementations restrict themselves to `syntax-rules` macros. Several more powerful `syntax` constructions over miniKanren may rely on more expressive macro systems, including Ballantyne’s `minikanren-ee` [1].

This note shows that at the cost of one additional feature—namely variadic functions—implementers provide a somewhat more powerful core language and significantly reduce the need for macros in implementing a language as expressive as the full microKanren + Macros approach provides.

Our desugaring macros would have worked whether we left or right associated. Removing a degree of freedom in the implementation, and forcing ourselves to program the solution functionally, led us to what seems like the right solution.

## References

- [1] Michael Ballantyne, Alexis King, and Matthias Felleisen. 2020. Macros for domain-specific languages. *Proceedings of the ACM on Programming Languages*, 4, OOPSLA, 1–29.
- [2] Daniel P. Friedman, William E. Byrd, Oleg Kiselyov, and Jason Hemann. 2018. *The Reasoned Schemer, Second Edition*. The MIT Press, (Mar. 2018). ISBN: 0-262-53551-3. [mitpress.mit.edu/books/reasoned-schemer-0](http://mitpress.mit.edu/books/reasoned-schemer-0).
- [3] Jason Hemann and Daniel P. Friedman. 2013. Mkanren: a minimal functional core for relational programming. In *Scheme 13*. <http://scheme-workshop.org/2013/papers/HemannMuKanren2013.pdf>.